



VIT[®]
Vellore Institute of Technology
(Deemed to be University under section 3 of UGC Act, 1956)

**SCHOOL OF COMPUTER SCIENCE ENGINEERING
AND INFORMATION SYSTEMS**

FALL SEMESTER 2024-2025

PMCA506L – CLOUD COMPUTING

DIGITAL ASSIGNMENT – 2

SUBMITTED ON: 23 – NOV - 2024

SUBMITTED BY-

AKASH KUMAR BANIK

PROGRAM: MCA

REGISTER No.: 24MCA0242

INSTALLATION AND CONFIGURATION OF JAVA 8 AND HADOOP 3.3.6 ON XUBUNTU (VIRTUALBOX)

STEP 1: INSTALLING JAVA 8

1.1 UPDATING THE SYSTEM

I began by ensuring my package manager was up to date. I opened the terminal in my xubuntu (light-weight version of ubuntu) virtual system and ran the command:

```
sudo apt update
```

This command refreshed the package repository information, allowing me to install the latest version of Java available in the repository.

1.2 INSTALLING JAVA

Next, I proceeded to install OpenJDK 8, which is the open-source implementation of the Java Platform. I executed the command:

```
sudo apt install openjdk-8-jdk
```

This command downloaded and installed the Java Development Kit (JDK) version 8. I could see the installation process in the terminal, where it fetched the required packages and set everything up.

When prompted, I pressed Y to allow the installation to proceed.

After the installation was complete, I wanted to confirm that Java was installed correctly. I checked the installed version by running the command:

```
java -version
```

The output displayed like this:

```
akb@akb-VirtualBox:~$ java -version
openjdk version "1.8.0_422"
OpenJDK Runtime Environment (build 1.8.0_422-8u422-b05-1~22.04-b05)
OpenJDK 64-Bit Server VM (build 25.422-b05, mixed mode)
akb@akb-VirtualBox:~$ S
```

This output indicated that Java 8 was successfully installed on my virtual system.

1.3 SETTING UP JAVA_HOME

Hadoop needs to know where Java is installed, so I had to set the JAVA_HOME environment variable. To find where Java was installed, I used:

```
dirname $(dirname $(readlink -f $(which java)))
```

This gave me the path /usr/lib/jvm/java-8-openjdk-amd64/jre.

```
akb@akb-VirtualBox:~$ dirname $(dirname $(readlink -f $(which java)))  
/usr/lib/jvm/java-8-openjdk-amd64/jre  
akb@akb-VirtualBox:~$
```

I then added this path to the .bashrc file. To do that, I opened .bashrc using:

```
nano ~/.bashrc
```

At the bottom of the file, I added:

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/jre  
export HADOOP_HOME=/home/akb/hadoop  
export HADOOP_INSTALL=$HADOOP_HOME
```

After saving the file, I reloaded the settings with the command:

```
source ~/.bashrc
```

```
akb@akb-VirtualBox:~$ source ~/.bashrc  
akb@akb-VirtualBox:~$
```

STEP 2: INSTALLING HADOOP 3.3.6

2.1 DOWNLOADING HADOOP

Next, I moved on to downloading Hadoop. I went to the Hadoop website and copied the download link for **Hadoop 3.3.6**. Then, in the terminal, I used wget to download hadoop:

```
wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz
```

```
akb@akb-VirtualBox:~$ wget https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz  
--2024-10-23 19:07:04-- https://downloads.apache.org/hadoop/common/hadoop-3.3.6/hadoop-3.3.6.tar.gz  
Resolving downloads.apache.org (downloads.apache.org)... 135.181.214.104, 88.99.208.237, 2a01:4f8:10a:39da::2, ...  
Connecting to downloads.apache.org (downloads.apache.org)|135.181.214.104|:443... connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 730107476 (696M) [application/x-gzip]  
Saving to: 'hadoop-3.3.6.tar.gz.2'  
  
hadoop-3.3.6.tar.gz.2 18%[====>] 129.20M 508KB/s eta 14m 45s
```

The download took some time, but after it was done, I unzipped the file using:

```
tar -xvzf hadoop-3.3.6.tar.gz
```

```
akb@akb-VirtualBox:~$  
akb@akb-VirtualBox:~$ tar -xvzf hadoop-3.3.6.tar.gz  
hadoop-3.3.6/  
hadoop-3.3.6/NOTICE-binary  
hadoop-3.3.6/licenses-binary/  
hadoop-3.3.6/licenses-binary/LICENSE-lz4.txt  
hadoop-3.3.6/licenses-binary/LICENSE-zstd-jni.txt
```

2.2 RENAMING THE HADOOP DIRECTORY

To make things easier to navigate, I renamed the folder from hadoop-3.3.6 to just hadoop:

```
mv hadoop-3.3.6 hadoop
```

2.3 SETTING UP HADOOP ENVIRONMENT VARIABLES

Now, I needed to set up environment variables for Hadoop. So, I edited the .bashrc file again using:

```
nano ~/.bashrc
```

At the bottom, I added the following lines to set up the Hadoop paths:

```
export HADOOP_HOME=~/.hadoop  
export HADOOP_INSTALL=$HADOOP_HOME  
export HADOOP_MAPRED_HOME=$HADOOP_HOME  
export HADOOP_COMMON_HOME=$HADOOP_HOME  
export HADOOP_HDFS_HOME=$HADOOP_HOME  
export HADOOP_YARN_HOME=$HADOOP_HOME  
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native  
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
```

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/jre  
export HADOOP_HOME=/home/akb/hadoop  
export HADOOP_INSTALL=$HADOOP_HOME  
export HADOOP_MAPRED_HOME=$HADOOP_HOME  
export HADOOP_COMMON_HOME=$HADOOP_HOME  
export HADOOP_HDFS_HOME=$HADOOP_HOME  
export HADOOP_YARN_HOME=$HADOOP_HOME  
export HADOOP_COMMON_LIB_NATIVE=$HADOOP_HOME/lib/native  
export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin  
export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native"
```

I saved the file and ran the command: `source ~/.bashrc`

This applied the new environment settings for Hadoop.

STEP 3: CONFIGURING HADOOP

Now that Hadoop was installed, I needed to configure several important files located in the `~/hadoop/etc/hadoop/` directory for the functioning of MapReduce properly.

3.1 CONFIGURING CORE-SITE.XML

First, I configured the **core-site.xml** file, which defines the default file system:

```
nano ~/hadoop/etc/hadoop/core-site.xml
```

Inside this file, I added:

```
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
```

```
<configuration>
<property>
  <name>fs.defaultFS</name>
  <value>hdfs://localhost:9000</value>
</property>
</configuration>
```

This ensures that Hadoop knows to use the localhost as the default file system for its operations.

3.2 CONFIGURING HDFS-SITE.XML

Next, I configured the **hdfs-site.xml** file to set up the HDFS (Hadoop Distributed File System) inside the Hadoop directory:

```
nano hdfs-site.xml
```

I added these properties:

```
<property>
  <name>dfs.replication</name>
  <value>1</value>
```

```

</property>
<property>
  <name>dfs.name.dir</name>
  <value>/home/akb/hadoop/data/namenode</value>
</property>
<property>
  <name>dfs.data.dir</name>
  <value>/home/akb/hadoop/data/datanode</value>
</property>

```

```

<configuration>
<property>
  <name>dfs.replication</name>
  <value>3</value>
</property><property>
  <name>dfs.name.dir</name>
  <value>/home/akb/hadoop/data/namenode</value>
</property><property>
  <name>dfs.data.dir</name>
  <value>/home/akb/hadoop/data/datanode</value>
</property>
</configuration>

```

I created the necessary directories for the NameNode and DataNode:

```
mkdir -p ~/hadoop/data/namenode
```

```
mkdir -p ~/hadoop/data/datanode
```

3.3 CONFIGURING MAPRED-SITE.XML

I then configured the MapReduce framework by configuring the mapred-site.xml file in the hadoop directory using the command:

```
nano mapred-site.xml
```

Inside, I added the properties:

```

<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>

```

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
<property>
```

3.4 Configuring yarn-site.xml

Finally, I configured **YARN**, Hadoop's resource manager inside the Hadoop directory using:

nano yarn-site.xml

I added these properties:

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
```

```
<configuration>
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property><property>
  <name>yarn.nodemanager.auxservices.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<!-- Site specific YARN configuration properties -->
</configuration>
```

STEP 4: SETTING UP SSH

Hadoop requires SSH to communicate between different nodes. I set up SSH on my virtual machine as follows:

4.1 GENERATING SSH KEY

I created an SSH key without a password by running:

ssh-keygen -t rsa

And pressed Enter on each prompted asked, until the ssh key is generated.

```
akb@akb-VirtualBox:~/hadoop/etc/hadoop$ ssh-keygen -t rsa
Generating public/private rsa key pair.
Enter file in which to save the key (/home/akb/.ssh/id_rsa):
/home/akb/.ssh/id_rsa already exists.
Overwrite (y/n)? y
Enter passphrase (empty for no passphrase):
Enter same passphrase again:
Your identification has been saved in /home/akb/.ssh/id_rsa
Your public key has been saved in /home/akb/.ssh/id_rsa.pub
The key fingerprint is:
SHA256:wavsMUzyz0+361Dx7BpTR10Vew9C4MXNrCGivCJwTFM akb@akb-VirtualBox
The key's randomart image is:
+---[RSA 3072]-----+
|  .E      .oo+ . = |
|  o      ....oo + + |
|  o . . .o...oo..+ |
|  . o   o o   . = oo |
|  o . ..S   . + o |
|  . *..   . o . |
|  . B   o + . |
|  . = . o = |
|  . o...= . |
+-----[SHA256]-----+
akb@akb-VirtualBox:~/hadoop/etc/hadoop$
```

4.2 AUTHORIZING SSH KEY

Then, I authorized the SSH key using the commands:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

```
chmod 640 ~/.ssh/authorized_keys
```

4.3 TESTING SSH

I tested my SSH setup by running the command: `ssh localhost`

```
akb@akb-VirtualBox:~$ cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
akb@akb-VirtualBox:~$ chmod 640 ~/.ssh/authorized_keys
akb@akb-VirtualBox:~$ ssh localhost
Welcome to Ubuntu 22.04.3 LTS (GNU/Linux 6.8.0-47-generic x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:       https://ubuntu.com/advantage

Expanded Security Maintenance for Applications is not enabled.

103 updates can be applied immediately.
To see these additional updates run: apt list --upgradable

20 additional security updates can be applied with ESM Apps.
Learn more about enabling ESM Apps service at https://ubuntu.com/esm

New release '24.04.1 LTS' available.
Run 'do-release-upgrade' to upgrade to it.

Last login: Wed Oct 23 19:54:01 2024 from 127.0.0.1
akb@akb-VirtualBox:~$
```


STEP 5: STARTING HADOOP

Now that everything was set up, I proceeded to start Hadoop.

5.1 FORMATTING NAMENODE

Before starting Hadoop, I needed to format the NameNode. I did this using the command:

hdfs namenode -format

5.2 STARTING HDFS

To start the Hadoop Distributed File System, I ran the command:

start-dfs.sh

5.3 STARTING YARN

To start the YARN resource manager, I ran the command:

start-yarn.sh

5.4 CHECKING THE SERVICES

To make sure everything was running, I used the jps command in the terminal:

jps

This command listed running Hadoop services like NameNode, DataNode, ResourceManager, and NodeManager.

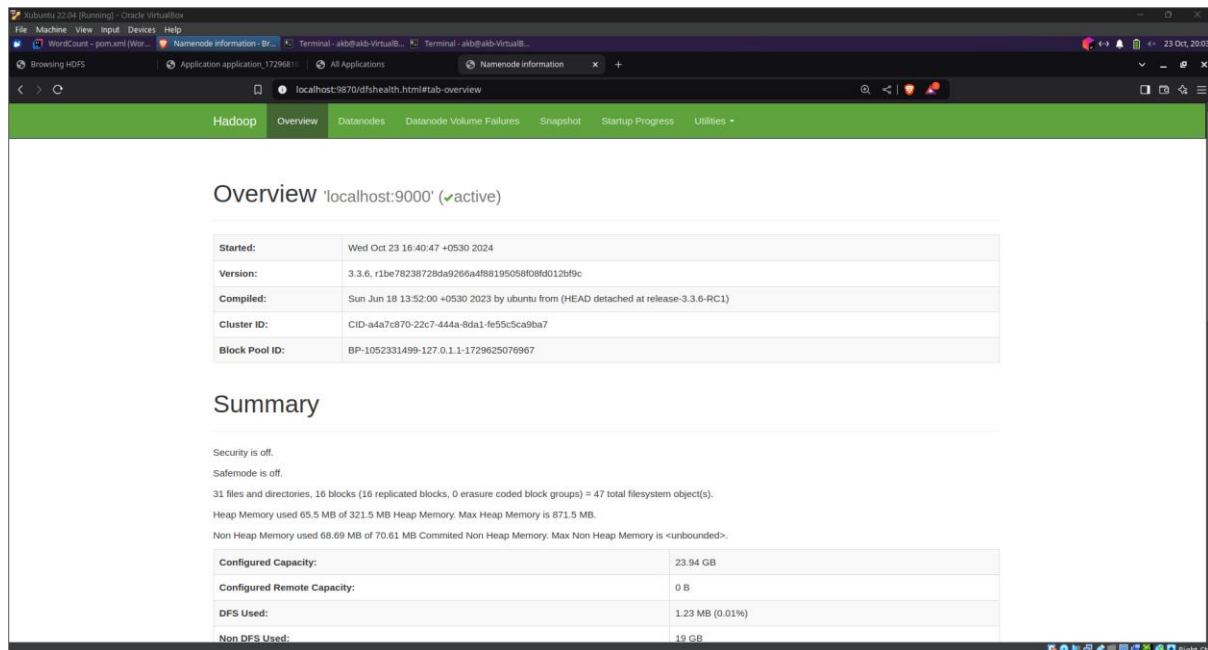
```
akb@akb-VirtualBox:~$ jps
6624 Jps
3412 SecondaryNameNode
2679 Main
3609 ResourceManager
2281
3066 NameNode
3726 NodeManager
3199 DataNode
akb@akb-VirtualBox:~$
```

STEP 6: ACCESSING HADOOP

Finally, to access the Hadoop web interface and check the status of my cluster, I opened the web browser and went to:

http://localhost:9870

This brought up the Hadoop cluster summary page, confirming that everything was working as expected.



ISSUES ENCOUNTERED AND SOLUTIONS

SSH Key Issue: Initially, I had permission issues with the SSH key. I fixed this by setting the right permissions:

chmod 640 ~/.ssh/authorized_keys

Connection Refused Error: When starting Hadoop, I encountered a "localhost connection refused" error. This was resolved by ensuring that SSH was installed and properly configured:

sudo apt install openssh-server

By following these steps, I was able to successfully install and configure Java 8 and Hadoop 3.3.6 on my Xubuntu virtual system, and now my system is ready to run Hadoop-based applications.

WORDCOUNT HADOOP PROJECT

In this project, I implemented a simple **WordCount** example using **Hadoop** and **Java**. The WordCount program is one of the most basic and widely used examples to demonstrate how Hadoop's MapReduce framework works. Its purpose is to count the occurrences of each word in a given text file by breaking down the process into two phases: **map** and **reduce**. This distributes the workload across multiple nodes for parallel processing, making it scalable for large datasets.

SETTING UP THE ENVIRONMENT

Before diving into the code, I needed to set up the environment, which involved installing and configuring **Maven**, **Hadoop**, and **IntelliJ IDEA**.

1. Setting Up IntelliJ IDEA and Maven

I used **IntelliJ IDEA Community Edition** to build this project. To start, I created a new **Maven** project.

- Opened IntelliJ and clicked on **New Project**.
- Selected **Maven** as the project type and ensured that **Java JDK** is installed.
- I named the project as WordCount, and set the group and artifact IDs. I set the groupId as **org.akb** and left the artifact ID as default.

After the project was created, I deleted the default main class, as it wasn't needed for this project.

The next step was to add the necessary **dependencies** to the pom.xml file.

2. Adding Dependencies

In the pom.xml file, I added dependencies for **Hadoop Common** and **Hadoop MapReduce Client Core** to ensure that my project could use the required Hadoop libraries.

```
<dependencies>
  <dependency>
    <groupId>org.apache.hadoop</groupId>
    <artifactId>hadoop-common</artifactId>
    <version>3.3.6</version>
```

```

</dependency>

<dependency>

    <groupId>org.apache.hadoop</groupId>

    <artifactId>hadoop-mapreduce-client-core</artifactId>

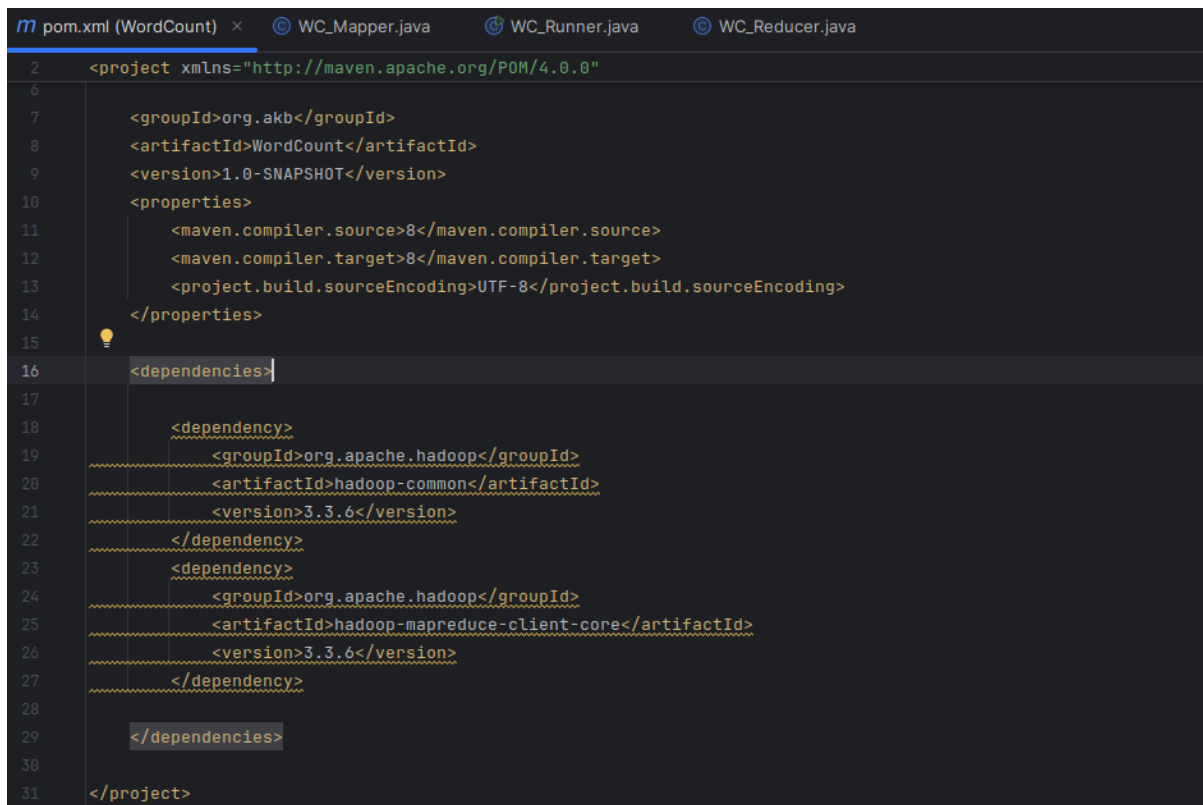
    <version>3.3.6</version>

</dependency>

</dependencies>

```

I reloaded the Maven project so that all dependencies were downloaded and added to the classpath.



```

m pom.xml (WordCount) x WC_Mapper.java WC_Runner.java WC_Reducer.java
2 <project xmlns="http://maven.apache.org/POM/4.0.0"
6
7 <groupId>org.akb</groupId>
8 <artifactId>WordCount</artifactId>
9 <version>1.0-SNAPSHOT</version>
10 <properties>
11 <maven.compiler.source>8</maven.compiler.source>
12 <maven.compiler.target>8</maven.compiler.target>
13 <project.build.sourceEncoding>UTF-8</project.build.sourceEncoding>
14 </properties>
15
16 <dependencies>
17
18 <dependency>
19 <groupId>org.apache.hadoop</groupId>
20 <artifactId>hadoop-common</artifactId>
21 <version>3.3.6</version>
22 </dependency>
23 <dependency>
24 <groupId>org.apache.hadoop</groupId>
25 <artifactId>hadoop-mapreduce-client-core</artifactId>
26 <version>3.3.6</version>
27 </dependency>
28
29 </dependencies>
30
31 </project>

```

WRITING THE WORDCOUNT PROJECT CODE

The WordCount project consists of three main components:

1. **WC_Mapper** – Handles the **map** phase of the process.
2. **WC_Reducer** – Handles the **reduce** phase of the process.
3. **WC_Runner** – Acts as the driver to set up and run the MapReduce job.

1. Creating WC_Mapper

I created a new class called WC_Mapper.java under the package org.akb. The mapper's job is to split the input text into individual words and emit each word as a key with a count of one as its value.

```
m pom.xml (WordCount)  WC_Mapper.java  WC_Runner.java  WC_Reducer.java
1 package org.akb;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5 import org.apache.hadoop.io.IntWritable;
6 import org.apache.hadoop.io.LongWritable;
7 import org.apache.hadoop.io.Text;
8 import org.apache.hadoop.mapred.MapReduceBase;
9 import org.apache.hadoop.mapred.Mapper;
10 import org.apache.hadoop.mapred.OutputCollector;
11 import org.apache.hadoop.mapred.Reporter;
12
13 public class WC_Mapper extends MapReduceBase implements Mapper<LongWritable,Text,Text,IntWritable> { 1 usage
14     private final static IntWritable one = new IntWritable( value: 1); 1 usage
15     private Text word = new Text(); 2 usages
16     @Override
17     public void map(LongWritable key, Text value,OutputCollector<Text,IntWritable> output, Reporter reporter) throws IOException{
18         String line = value.toString();
19         StringTokenizer tokenizer = new StringTokenizer(line);
20         while (tokenizer.hasMoreTokens()){
21             word.set(tokenizer.nextToken());
22             output.collect(word, one);
23         }
24     }
25 }
```

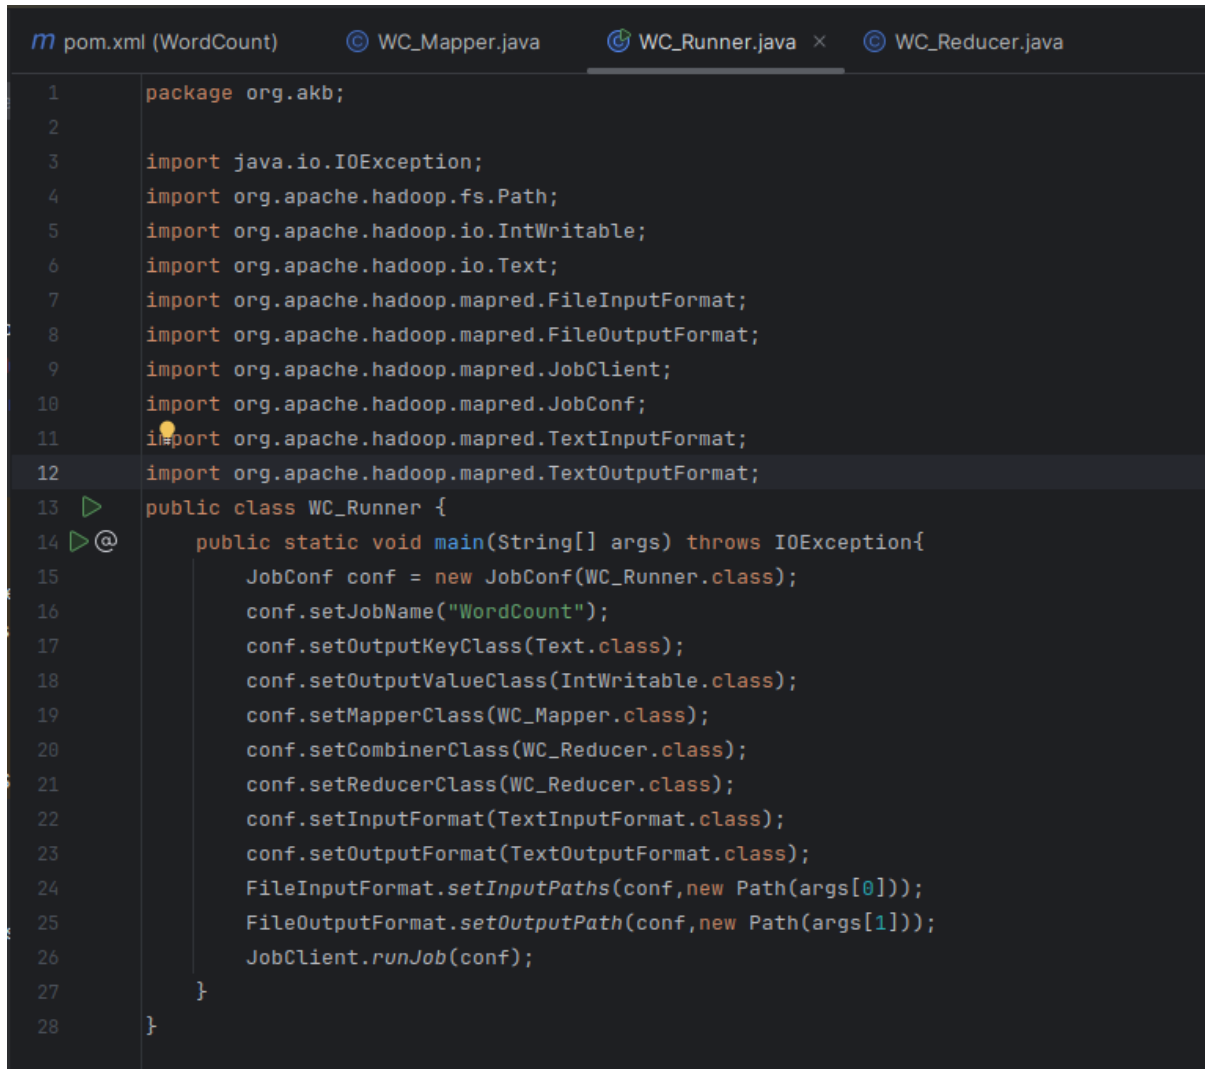
2. Creating WC_Reducer

Next, I created WC_Reducer.java. The reducer receives each word and the corresponding list of counts from the mapper. It sums the counts for each word and outputs the final result.

```
m pom.xml (WordCount)  WC_Mapper.java  WC_Runner.java  WC_Reducer.java x
1 package org.akb;
2
3 import java.io.IOException;
4 import java.util.Iterator;
5 import org.apache.hadoop.io.IntWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapred.MapReduceBase;
8 import org.apache.hadoop.mapred.OutputCollector;
9 import org.apache.hadoop.mapred.Reducer;
10 import org.apache.hadoop.mapred.Reporter;
11
12 public class WC_Reducer extends MapReduceBase implements Reducer<Text,IntWritable,Text,IntWritable> { 2 usages
13     @Override
14     public void reduce(Text key, Iterator<IntWritable> values,OutputCollector<Text,IntWritable> output,
15         Reporter reporter) throws IOException {
16         int sum=0;
17         while (values.hasNext()) {
18             sum+=values.next().get();
19         }
20         output.collect(key,new IntWritable(sum));
21     }
22 }
```

3. Creating WC_Runner

Finally, I created WC_Runner.java, which configures and runs the MapReduce job. This class defines the input and output locations and specifies the Mapper and Reducer classes.



```
1 package org.akb;
2
3 import java.io.IOException;
4 import org.apache.hadoop.fs.Path;
5 import org.apache.hadoop.io.IntWritable;
6 import org.apache.hadoop.io.Text;
7 import org.apache.hadoop.mapred.FileInputFormat;
8 import org.apache.hadoop.mapred.FileOutputFormat;
9 import org.apache.hadoop.mapred.JobClient;
10 import org.apache.hadoop.mapred.JobConf;
11 import org.apache.hadoop.mapred.TextInputFormat;
12 import org.apache.hadoop.mapred.TextOutputFormat;
13
14 public class WC_Runner {
15     public static void main(String[] args) throws IOException{
16         JobConf conf = new JobConf(WC_Runner.class);
17         conf.setJobName("WordCount");
18         conf.setOutputKeyClass(Text.class);
19         conf.setOutputValueClass(IntWritable.class);
20         conf.setMapperClass(WC_Mapper.class);
21         conf.setCombinerClass(WC_Reducer.class);
22         conf.setReducerClass(WC_Reducer.class);
23         conf.setInputFormat(TextInputFormat.class);
24         conf.setOutputFormat(TextOutputFormat.class);
25         FileInputFormat.setInputPaths(conf, new Path(args[0]));
26         FileOutputFormat.setOutputPath(conf, new Path(args[1]));
27         JobClient.runJob(conf);
28     }
29 }
```

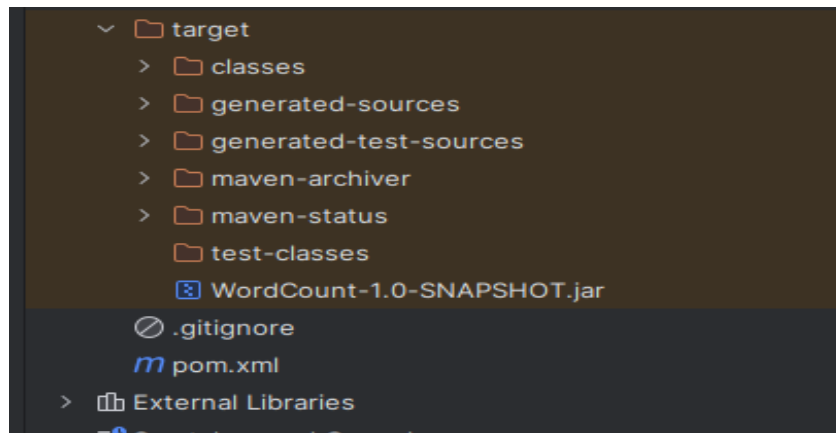
BUILDING AND RUNNING THE PROJECT

1. Creating the JAR File

Once all the code was written, I built the project using Maven. I opened the terminal in IntelliJ and ran the following commands to clean and package the project:

mvn clean package

This generated a JAR file in the target folder, which I would use to run the WordCount job.



2. Running the WordCount Job

Next, I created a simple text file as input. I used the terminal to create a file called input.txt with some random text using the commands:

```
nano input.txt
```

I added the following text:

This is the input file for hadoop project

This is for Cloud Computing with AKB hadoop project file

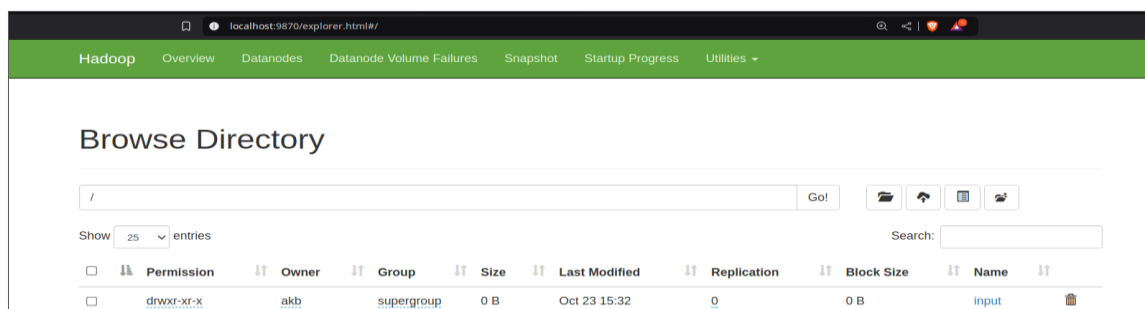
This is done for now

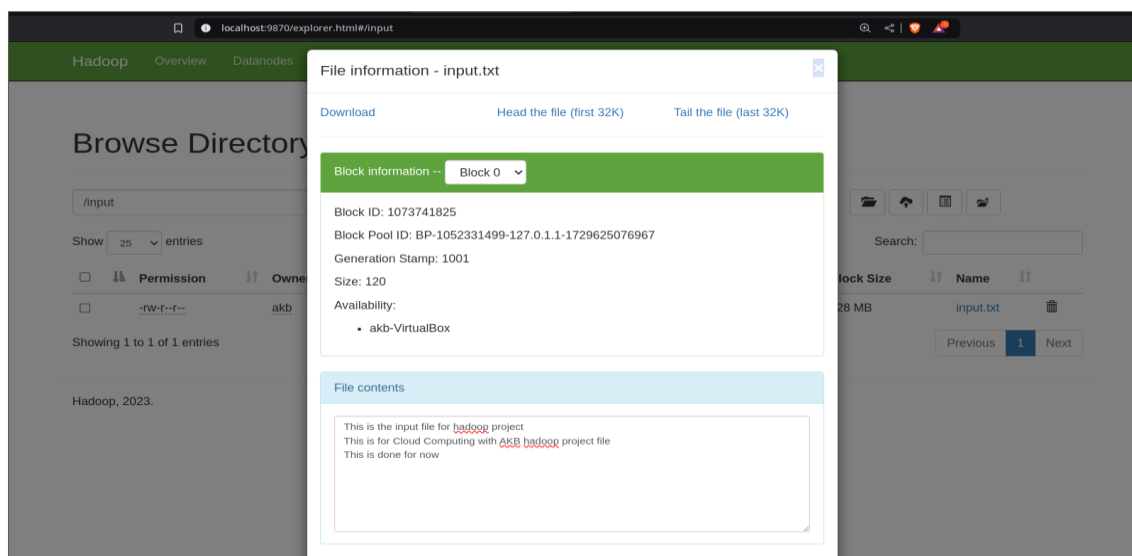
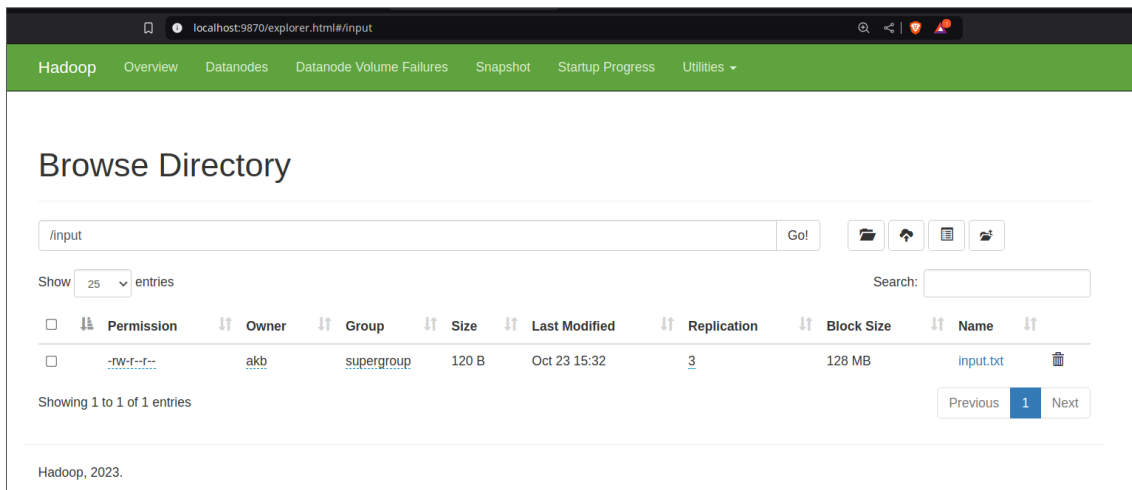
I then created a directory in HDFS to store the input file:

```
hadoop fs -mkdir /input
```

Then, I copied the input file to HDFS:

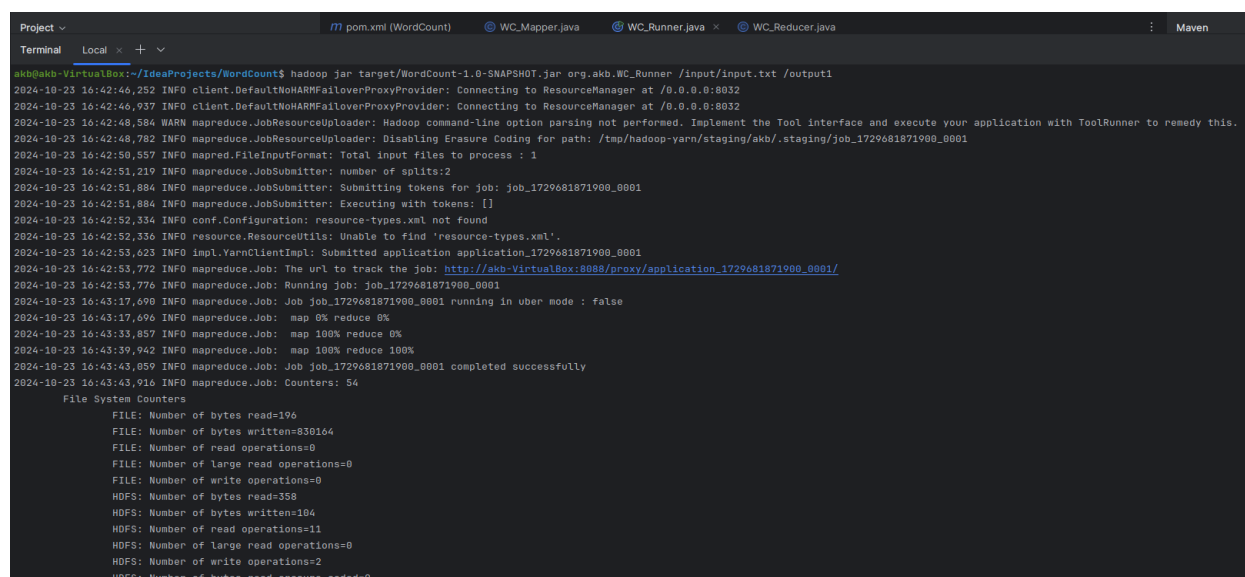
```
hadoop fs -put input.txt /input
```





Finally, I ran the WordCount job by executing the following command:

hadoop jar target/wordcount-1.0-SNAPSHOT.jar org.akb.WC_Runner /input/input.txt /output




```

Job Counters
  Launched map tasks=2
  Launched reduce tasks=1
  Data-local map tasks=2
  Total time spent by all maps in occupied slots (ms)=27925
  Total time spent by all reduces in occupied slots (ms)=3539
  Total time spent by all map tasks (ms)=27925
  Total time spent by all reduce tasks (ms)=3539
  Total vcore-milliseconds taken by all map tasks=27925
  Total vcore-milliseconds taken by all reduce tasks=3539
  Total megabyte-milliseconds taken by all map tasks=28595200
  Total megabyte-milliseconds taken by all reduce tasks=3623936
Map-Reduce Framework
  Map input records=3
  Map output records=23
  Map output bytes=212
  Map output materialized bytes=202
  Input split bytes=178
  Combine input records=23
  Combine output records=17
  Reduce input groups=14
  Reduce shuffle bytes=202
  Reduce input records=17
  Reduce output records=14
  Spilled Records=34
  Shuffled Maps =2
  Failed Shuffles=0
  Merged Map outputs=2
  GC time elapsed (ms)=1917
  CPU time spent (ms)=3330
  Physical memory (bytes) snapshot=749191168
  Virtual memory (bytes) snapshot=7643070464
  Total committed heap usage (bytes)=740818944
  Peak Map Physical memory (bytes)=253992960
  Peak Map Virtual memory (bytes)=2544386048
  Peak Reduce Physical memory (bytes)=241606656
  Peak Reduce Virtual memory (bytes)=2554531840
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=180
File Output Format Counters
  Bytes Written=104
akb@akb-VirtualBox:~/IdeaProjects/WordCount$

```

3. Checking the Output

Once the job finished, I verified the output by listing the contents of the /output directory in HDFS:

```
hadoop fs -cat /output/part-00000
```

The output displayed the word count:

```

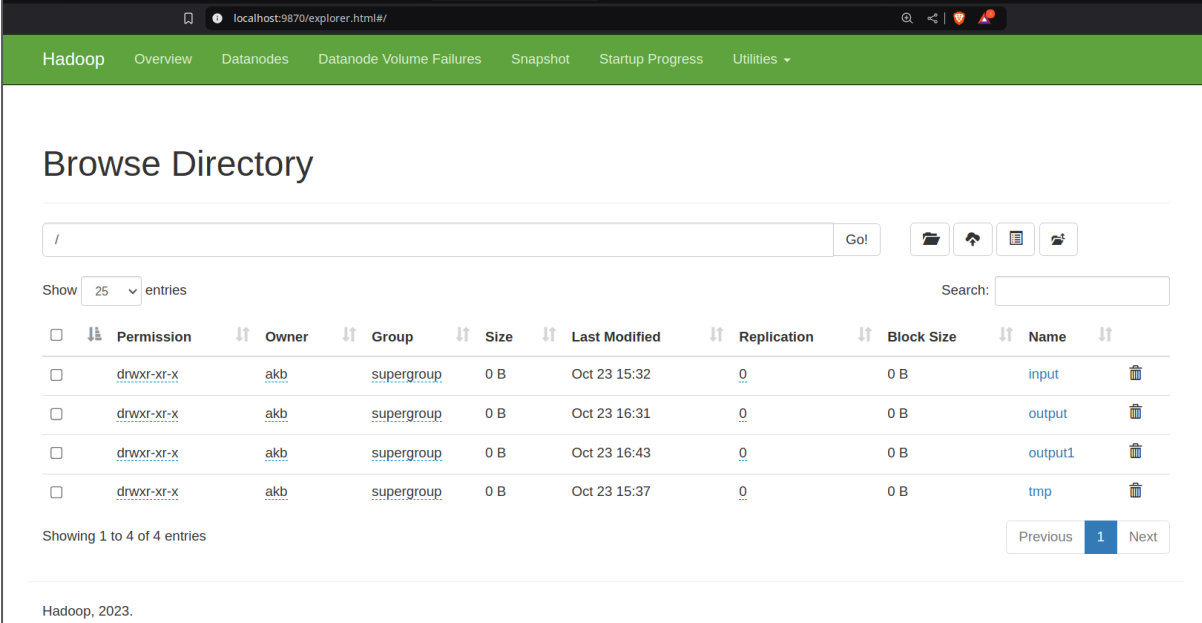
akb@akb-VirtualBox:~$ hadoop fs -cat /output/part-00000
AKB      1
Cloud    1
Computing      1
This      3
done      1
file      2
for      3
hadoop    2
input     1
is        3
now       1
project   2
the       1
with      1
akb@akb-VirtualBox:~$

```

CONCLUSION

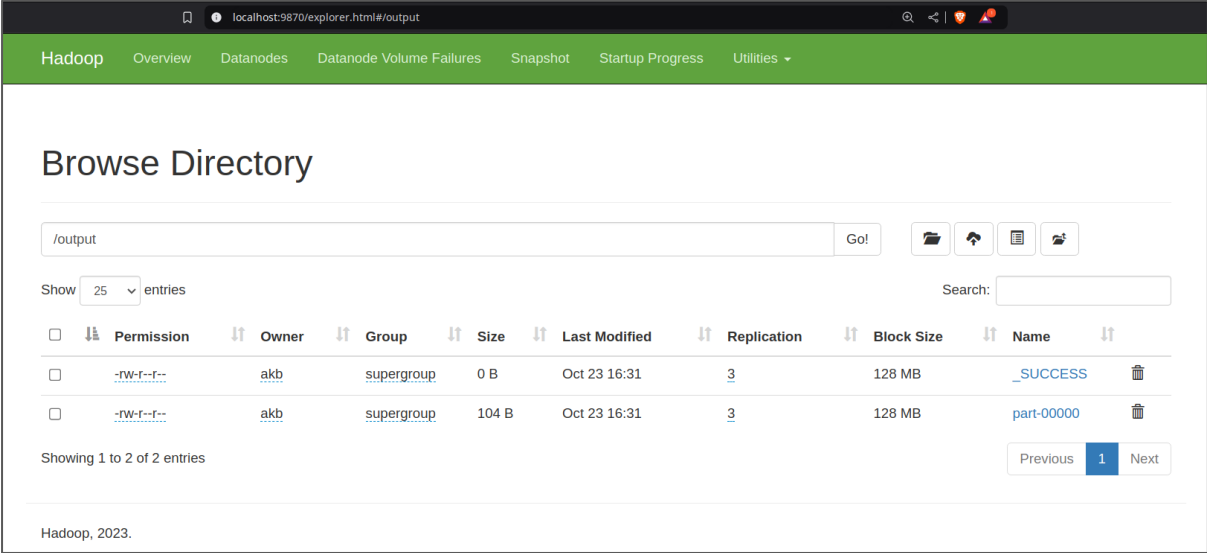
By following these steps, I successfully created and ran a WordCount Hadoop project using Maven and IntelliJ IDEA. The WordCount example demonstrates how Hadoop's MapReduce framework works by distributing the workload and processing large amounts of data in parallel.

OUTPUT SCREENSHOTS



The screenshot shows the Hadoop File Explorer interface at the root directory. The breadcrumb is "/". The table lists four files: input, output, output1, and tmp, all with permissions drwxr-xr-x, owned by akb, in the supergroup, with a size of 0 B and a replication factor of 0. The last modified times are Oct 23 15:32, Oct 23 16:31, Oct 23 16:43, and Oct 23 15:37 respectively. The interface shows "Showing 1 to 4 of 4 entries" and navigation buttons for Previous, 1, and Next.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
drwxr-xr-x	akb	supergroup	0 B	Oct 23 15:32	0	0 B	input
drwxr-xr-x	akb	supergroup	0 B	Oct 23 16:31	0	0 B	output
drwxr-xr-x	akb	supergroup	0 B	Oct 23 16:43	0	0 B	output1
drwxr-xr-x	akb	supergroup	0 B	Oct 23 15:37	0	0 B	tmp



The screenshot shows the Hadoop File Explorer interface at the /output directory. The breadcrumb is "/output". The table lists two files: _SUCCESS and part-00000, both with permissions -rw-r--r--, owned by akb, in the supergroup, with a size of 0 B and a replication factor of 3. The last modified times are both Oct 23 16:31. The _SUCCESS file has a block size of 128 MB. The interface shows "Showing 1 to 2 of 2 entries" and navigation buttons for Previous, 1, and Next.

Permission	Owner	Group	Size	Last Modified	Replication	Block Size	Name
-rw-r--r--	akb	supergroup	0 B	Oct 23 16:31	3	128 MB	_SUCCESS
-rw-r--r--	akb	supergroup	104 B	Oct 23 16:31	3	128 MB	part-00000

The screenshot displays the Hadoop web interface at localhost:9870/explorer.html#/output. The main panel shows the 'Browse Directory' view for the /output path, listing two entries with permissions -rw-r--r-- and owner akb. A modal window is open, showing 'Block information -- Block 0'. The block details are as follows:

- Block ID: 1073741846
- Block Pool ID: BP-1052331499-127.0.1.1-1729625076967
- Generation Stamp: 1022
- Size: 104
- Availability:
 - akb-VirtualBox

The 'File contents' section shows the following text:

```
AKB 1
Cloud 1
Computing 1
This 3
done 1
file 2
for 3
hadoop 2
```

The background interface also shows a table with columns 'k Size' and 'Name', listing files like _SUCCESS and part-00000.